Lazada Product Title Quality Challenge

Rodrigo Valle (104494120)   Bill Tang (104621566)
Krit Saefang (904723127)

March 16, 2018

Abstract

In this paper, we search for a reasonable approach to classify titles as either clear or unclear and concise or not concise. This might be useful for websites with user submitted content, who would want to flag titles which do a poor job of communicating their intent in order to improve user experience. Being significantly cheaper and less tedious than paying someone to manually label the quality of a title, especially with a reasonable error rate, we believe our solution is valuable. We explore, build, and evaluate three major machine learning models for classifying titles: Naive Bayes, Support Vector Machines (SVMs), and Neural Networks, as well as multiple ways to summarize titles into feature vectors to be used in conjunction with our models.

Background and Problem Definition

Title classification is done on a dataset with product titles and their associated information including: description, categories and price, along with a few less relevant tags. Each title in the training set has been labeled individually by Lazada's quality control team as either clear or unclear and concise or not concise. Our goal is to train a model that, given a product title and its associated information, can predict the probability that the title belongs to each of the "clear" and "concise" classes. Models will be evaluated by using the Root Mean Squared Error (RMSE) of the models predictions when compared to the ground-truth labels, defined as follows:

$$\text{RMSE} = \sqrt{\sum \frac{(y_{\text{pred}} - y_{\text{ref}})^2}{N}}$$

Since labels for the test and validation data on the competition website are not provided, we will be splitting the test set of roughly thirty-six thousand titles into test and training sets, with sizes of about 4/5 and 1/5 of the original data respectively. This should still be more than enough data to train our classifiers, some of which have long runtimes, although the additional data might improve our results. This is something we discuss in the future work section of the SVM model.

Title classification can be tricky for various reasons, the most concerning one being that labels such as "clear" and "concise" are not well defined. The dataset was labeled by several different people, so there may be even less concistency between examples. In general, the dataset will have noise and coupled with the fact that these concepts are not well defined and subject to personal preference, a perfect classifier is simply impossible to achieve. However, we can still do a reasonably good job of classifying titles regardless.

Methods

Naïve Bayes

Introduction

Naïve Bayes are a family of probabilistic classifiers based on Bayes' Theorem. These classifiers are simple and are applicable to many uses, such as document classification and spam filtering. Compared with other methods studied in this report, they are fast to train and, in theory, does not require much data to train on. Compared to neural networks, Naïve Bayes is a calculation, and thus less demanding in processing power and resource requirements.

For the project, three implementations of the Naïve Bayes were used. These are of the scikit-learn package as is the SVM implementation used. The first is Gaussian Naïve Bayes, an implementation where the likelihood of the features is assumed to be Gaussian. Second is the Multinomial Naïve Bayes, where the data us assumed to follow the multinomial distribution. The multinomial implementation is one of the classic variants used for text classifications. The third and final implementation is the Bernoulli Naïve Bayes, suited for data with multivariate Bernoulli distributions, and would work well with binary valued data. [1]

The given test data was split up at the 11000 from last item mark, roughly two thirds of the given data. Everything up to the last 11000 entries are used as training data, given to the classifiers along with the label, and the remaining 11000 is used for validation testing, using the labels to check for errors. The data will be vectorized before being fed into the classifiers.

All three implementations were fed the same data, features, and, if applicable, tested with same settings, such as priors and smoothing settings. Tests were done with base settings (Title or description only, with no further tweaking), without smoothing applied, using titles along with categories, using titles along with categories but without smoothing, without using fit priors, and using with class priors. Class priors were given as [0.685,0.315] and [0.943,0.053] for clarity and concise respectively, calculated directly from the dataset. In addition, another feature list for testing conciseness, shared with the neural network tests, will be used as well.

Gaussian Naïve Bayes

The Gaussian Naïve Bayes implementation was the first to be used. The implementation allows feeding in the prior probabilities of the class. Compared to the other Naïve Bayes implementations tests, Gaussian Naïve Bayes suffers from higher resource usage and is much

slower. This is a consequence of requiring matrices be dense matrices. In order to get Gaussian Naïve Bayes to actually run, the virtual machine used to do the tests had to be given about 10GB of memory to work with, and even with this, applying Gaussian Naïve Bayes for the clarity tests failed due to not enough allocated memory. Because data batching was not done, results for Gaussian Naïve Bayes were skipped for the clarity tests. [2]

The results of conciseness test were poor compared to the other implementations. The RMSE for the base fit was 0.690, and remained approximately the same across the tests using both titles and categories, and with class priors. Using the features shared with the neural network tests, the results improved to 0.577.

	RMSE for GaussianNB
Base	0.690
Title with categories	0.690
With class priors	0.690
Generated features	0.577

Multinomial Naïve Bayes

The implementation for Multinomial Naïve Bayes performed much better than the Gaussian implementation. To start, it did not require matrices being fed in to be dense, and as such sparse matrices works. Unlike the Gaussian implementation, Multinomial Naïve Bayes actually has a smoothing function, and allows the toggling of fit priors. [3]

Because of the sparse matrices, training and predicting are also much faster. And even more, it was able to complete the clarity tests as well. For the concise tests, base results without medications resulted in the best RMSE results at 0.460. Multinomial Naïve Bayes performed very well with the clarity tests, clocking 0.275 as its best RMSE result, on base settings. Turning off smoothing and adding features does not seem to help in terms of improving results, nor did disabling fit priors and providing class priors. The generated features used to train our neural-network implementation do not work with Multinomial Naïve Bayes due to having non-discrete negative values and as such were skipped.

	RMSE for MultinomialNB - Concise	RMSE for MultinomialNB - Clarity
Base	0.460	0.275
Without smoothing	0.491	0.362
Titles with categories	0.496	0.311
Titles with categories without smoothing	0.508	0.447
Without fit priors	0.478	0.320
Without class priors	0.501	0.401
Generated features	No data	

Bernoulli Naïve Bayes

The last Naïve Bayes implementation tested was the Bernoulli Naïve Bayes. Like Gaussian Naïve Bayes, Bernoulli Naïve Bayes ran into issues with exceeding available memory, although this time it was solved by using a hashing vectorizer instead of the count vectorizer as used by both Gaussian and Multinomial implementations. Training and prediction speed was faster than Gaussian, and not too much different from Multinomial. Like the Multinomial implementation, it has a smoothing function, and allows the toggling of fit priors in addition to being able to use class priors. [4]

Like the Multinomial Naïve Bayes implementatioe, Bernoulli Naïve Bayes didn't result in any significant RMSE results for the clarity tests, but did do well with the concise test. On average, the results for the clarity tests were higher than those provided by Multinomial, with the best result coming from turning off smoothing. For the clarity test, Bernoulli Naïve Bayes achieved the best results out of the three implementations scoring 0.237 at the lowest using only the title and title with categories.

	RMSE for BernoulliNB - Concise	RMSE for BernoulliNB - Clarity
Base	0.560	0.237
Without smoothing	0.479	0.384
Titles with categories	0.560	0.237
Titles with categories without smoothing	0.508	0.401
Without fit priors	0.560	0.751
Without class priors	0.560	0.751
Generated features	0.528	

Conclusions with Naïve Bayes

While Naïve Bayes, compared to the other models in this report, did not give good results for clarity, they did decently well with the concise label. These numbers are, however, still unmatched by results from our tests with the Support Vector Machine (SVM) and Deep Neural Network (DNN), both of which provided much better results in the sub 0.1s mean-squared error. DNN in particular, reached an MSE result of 0.045. However, compared to Tensorflow as will be shown, Naïve Bayes can give its results with significantly less resources required and time spent.

An interesting observation is that under almost all circumstances, using the base settings alone resulted in some of the best predictions available from Naïve Bayes, and that adding priors or other features such as categories into the training only serves to worsen the results. Parts of this could be argued that the requirements for scoring these tests didn't really look into the accuracy of the titles and descriptions themselves, and other attributes such as length may play bigger role in being a good determinant for predictions.

The rather high results from the tests could also indicate that Naïve Bayes may not be the most suitable method for the given task, lacking in concepts that a neural network may be able to learn and adapt to. All the tests done with Naïve Bayes had training and

prediction done within or around one minute, and within the 10GB memory cap (with the exception of Gaussian Naïve Bayes which exceeded 10GB for the concise tests). Compare this to Tensorflow's 20 minute training time and reaching 40GB in data requirements and the trade-offs are clear.

Future work with Naïve Bayes

There are still more feature combinations that could be tested to see if results can be further improved. What was tested in this report is nowhere near exhaustive and other derivable features such as text length and combinations of titles and descriptions can be looked into. Different feature lists can be generated and used to as well to see the results of various combinations. While it can be argued that the method itself is not suitable for the given tasks, it doesn't mean it's not possible to improve the results, as even if the prediction is not as good as other methods, its speed and resource requirements can still lend it an upper hand under conditions where speed is required and resources lacking.

Support Vector Machine

The support vector machine classifier, at its core, separates labeled data in n dimensional space with an $n-1$ dimensional hyperplane. The best hyperplane is chosen as one that separates the two classes with the widest possible margin, measured as the closest member to the hyperplane from each class.

Most of the time though, data is not linearly seperable in n dimensions and so to get around this (and make support vector machines the stable and widely used tool that they are today), support vector machines will map n dimensional vectors into some higher dimensional space where the data *can* be separated by a hyperplane. [8]

Model Selection

We did our training of the SVM model using SciKit Learn's implementation, which in turn relies on the LibSVM library. There's a few different kernels that SciKit Learn lets us use for this higher dimensional mapping trick, but it's been the author's experience that of the four options: "linear", "polynomial", "rbf", and "sigmoid", the "rbf" kernel usually performs the best. This kernel choice was made in order to minimize the search space for good model hyperparameters in the later hyperparameter grid-search.

The model, with the Radial Basis Function (RBF) kernel specified, has two major tunable hyperparameters: C and gamma. The C parameter is a tradeoff between "misclassification of training examples agains the simplicity of the decision surface" [5]. A high C might classify more training examples correnctly at the expense of a more complicated separation, and a low C will cause the model to output a smoother separation surface at the expense of misclassifying more training examples. The gamma parameter has to do with the radial basis function kernel and the influence of each training example as related to distance. A higher value of gamma means a smaller sphere of influence for each training example, and could make the final model suceptible to overfitting. Likewise, a lower value of gamma

might combat overfitting, but we might lose some structural patterns in the original data, contributing to worse predictions.

Feature Generation from Titles

Arguably the most important part of a good support vector machine based model is a good feature space, where the different lables can be linearly separated. To get off the ground, we relied on SciKit Learn's CountVectorizer class, with the binary option set to true, to turn each product's title into a vector. Each dimension in this feature space represents a unique word. For a given title, a dimension is set to 1 if the corresponding unique word exists in the title, and 0 otherwise.

We chose this mapping of titles to features because alternative popular mappings such as term frequency-inverse document frequency or TD-IDF generate noisy feature vectors when used with a corpus of short strings such as titles, and its behavior is not suitable for our needs. TD-IDF works well on larger documents because it values words that appear frequently within a particular document (or small subset of documents), but penalizes words that appear frequently in most documents. This is not what we want because an overabundance of these globally frequently occurring words in any one particular title might signify a title of poor quality, and in addition, produces noisy feature vectors with short documents. The binary features that we chose to start with are simple, lightweight, and resistant to noise, even if they are high dimensional. [10]

Later in our investigation, it became apparent that better feature vectors were necessary if the model was going to improve. These binary word vectors did not capture enough of a title's semantic properties to allow the SVM to find a good separation.

To get around this, we employed the use of the Word2Vec based Doc2Vec algorithm to encode titles as vectors [9]. The Doc2Vec neural network trains on all titles in the training set in order to build a feature space where related documents are close to each other. For our purposes, we tagged each document with the three categories of product that it belonged to, in order to give Doc2Vec a better chance at guessing the semantic meaning behind a particular title. Then we can use the model to infer the feature vector of titles that we haven't trained on, and we can add a few other title measurements such as the length of the title in ASCII characters and also words, the number of nonalphanumeric characters present in the title, etc.

Grid Search for Hyperparameters

In order to choose reasonable values for the hyperparameters C and gamma, we performed a grid search across the ranges C $= [10^{-5}, 10^{-3}, ..., 10^{-13}, 10^{-15}]$ and gamma $= [10^{-15}, 10^{-13}, ..., 10^{1}]$ [7].

Gridsearch works by training an SVM on each pair of these parameters and performing cross validation to see which parameters yield an SVM with the lowest error. Gridsearch is expensive, so we had to be smart about our cross validation strategy.

Initally, we take the dataset and perform a stratified k-folds separation, where we split the dataset into 5 different sections and keep 4 for training and 1 for testing the completed model. We want the split to be "stratified" so that it keeps the same distribution of labels.

Our original dataset for clarity and concise labels is very skewed, with these labels being true with a high probability, so this is important for a fair evaluation of future performance.

Then, of the initial training set, we create our features as described above and perform a Stratified Shuffle Split, randomly selecting 1/5 of the input data to train the parametrized SVM and another 1/5 to evaluate it, while maintaining label distribution. This produces results quickly because we're not training on as much data. Combined with GridSearchCV's concurrency feature and we get cross validation results that look like the figures (1) and (2) in a reasonable amount of time (a few hours).

The results for the clarity gridsearch (figure (2)) were not so clear, so we ran it again but this time using the imblearn pacakge's SMOTE implementation with default settings. Adding SMOTE to our cross validation pipeline will redistribute the label frequencies by oversampling the minority class and undersampling the majority class. This allows us to see how well our model might perform on a more balanced dataset, which is useful for our gridsearch where we want to choose good hyperparameters. In general, this skew of our dataset can be used to our advantage, so we wouldn't want to train using SMOTE because it might cause the SVM to think that certain labels are more frequent than they actually are [6]. You can see our revised gridsearch results in figure (3).

Final Results

With the gridsearch results, we can see that the best value for gamma is approximately 10^{-1} and C is approximately 10^5, for both models. Using these values, we trained an SVM for each "concise" and "clarity" class using our binary word features and our Doc2Vec features. Our results are as follows:

Features	RMSE - Concise	RMSE - Clarity
Binary Word Vectors (no gridsearch)	0.4346	0.371
Binary Word Vectors (gridsearch)	0.3757	0.225
Doc2Vec Title Vectors (gridsearch)	0.4543	0.230

The untuned SVM trained on Binary Word vectors was trained on 26282 training examples and the parameters C = 10^3, gamma = 1/26282.

The RMSEs for the SVM trained after gridsearch were taken directly from the gridsearch's cross validation. The parameters for concise SVM were C = 10, gamma = 0.1, and for the clarity SVM they were C = 0.1 and gamma = 0.1.

The SVM trained on Doc2Vec title vectors with clarity labels was trained on only 1/5 of the labeled dataset, and tested on a separate 1/5 of the labeled dataset because training times were too extravagant. The performance of the SVM with doc2vec features seemingly decreased on the "concise" label; this could be due to a poor choice of hyperparameters. GridSearchCV could not be performed with doc2vec features because the computational power necessary was simply unavailable; perhaps a revision of the GridSearchCV process could speed things up further, or one of the alternate SVM models we describe below.

Future Work

Figure 1: Conciseness SVM Classifier Gridsearch

Figure 2: Clarity SVM Classifier Gridsearch

Figure 3: Clarity SVM Classifier Gridsearch with SMOTE

In the future, with such a large dataset, we ran into a lot of issues regarding SVM runtime. There are a couple ways to fix this; we could try using a LinearSVM whose training time is significantly shorter than the standard RBF SVM. There are several "linear learner" SVM implementations too, better suited for large datasets. We could also try training an ensemble of SVMs so as to be able to parallelize SVM training process.

The dataset also provides a lot of unlabeled validation and test data files, which our current models can't make use of. To get around this, we might try training a transductive SVM, or TSVM, which is a semilearning SVM model that can use unlabeled data to try and lower our RMSE scores. It would probably work to improve our testing, but runtime is too significant for even our standard SVM so this might prove to be too unwieldy to be useable.

A further improvement might be to use the label output from the clarity model in the conciseness model or vice versa, because looking at the data it's pretty clear that these two labels are somewhat dependent on each other.

Deep Neural Networks

Deep neural networks are powerful universal classifiers that have the ability to easily learn non-linear correlations between inputs and outputs. Relying on iterative training steps aimed at minimizing prediction loss, neural networks are able to gradually learn connections between features and labels, making them very effective models for classification. The main advantage of deep neural networks over other iterative approaches lie in the hidden layers that allow for the learning of arbitrary functions.

Feature extraction and selection

One of the major challenges faced in constructing the model came from feature selection. Our task was to classify title clarity and conciseness for products, and it was not immediately obvious which features should be selected, and how they should be transformed into numerical values that could be fed into a neural network.

First, our input data was preprocessed, using BeautifulSoup to remove unwanted html tags and cleaned so that words were all lowercase. The preprocessed features are combined with the clarity and conciseness labels are exported into a csv so they can be easily loaded and reutilized.

Initially, experimentation was conducted using simple easily extracted or calculated statistical features such as title length, number of title words, number of non alphanumeric characters, and other potentially related variables. Training on these initial features gave us RMSE values of around 0.48 for conciseness and 0.25 for clarity.

The vectorization of the groups of words formed the core of the feature set for the deep neural network model. After experimenting with bag of word models using tensorflow built in vocab processor with only marginal improvements in RMSE scores, we chose to use the Stanford's GloVe word embedding model in order to generate word embeddings that retained information about linguistic and semantic similarity [14]. To transform titles, categories, and description strings into vectorized features, the mean GloVe embedding vector was calculated by taking the mean of all the 200 dimensional word vectors generated from the

string words. This embedding allowed us to further generate interesting features like the relatedness between the title and description and categories by calculating the norm between the corresponding mean embedding vectors.

Finally, we also utilized linguistic features from the Spacy natural language processing libraries 'en' language model to generate more numeric features from our title, category, and description strings [15]. We extracted noun chunks from title strings, category strings, and description strings, and once again calculated mean embedding vector norms between them. We also incorporated counts of part of speech classification for words in the title strings. Put together, there are a total of 232 features that are fed into the input layer of the model.

Model Description

In this project, we built a feed-forward deep neural network on top of the Tensorflow framework using Keras. The model consists of the input layer, two dense hidden layers with hidden units of 100 and 50 units using ReLU activation functions, an output softmax layer with the two possible classes, and an additional dropout layer with a dropout rate of 0.2.

The ReLU or rectified linear unit is an activation function that has become very popular in recent years. It computes the function f(x) = max(0,x). It has been found to greatly improve the convergence speed of gradient descent optimizers when compared to sigmoid or tanh functions due to its linear and non-saturating form [13].

The choices for the hidden units and layers came from experimentation, although in terms of the final RMSE scores, the shape of the internal model did not change the results much as long as there were enough hidden units to avoid underfitting. We noted that increasing model depth past two hidden layers did not improve the model's performance, but slowed down the training significantly. Overfitting is protected against in our model through regularization in our dropout layer [16].
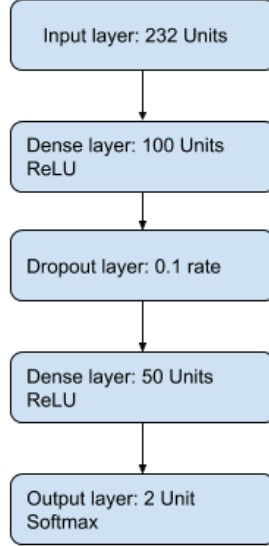
Figure 1: Model structure

Determining optimal training hyperparameters was also a challenging experimental process. After testing with various optimizers, we found the best results using an Adagrad optimizer with a learning rate of 0.01 and batch size of 128. At higher batch sizes convergence took many more epochs, while lower batch sizes made the training time for each epoch much higher. Other optimizers we considered were stochastic gradient descent, which performed poorly both in terms of convergence speed and model accuracy, and the Adam optimizer, which converged the fastest but with slightly worse RMSE scores. The Adagrad optimizer is a subgradient method that attempts to dynamically incorporate knowledge about the geometry of data observed in previous iterations [12]. Mean squared error was selected as the loss function since RMSE was the final model metric that we wanted to optimize. In terms of training epochs, 200 epochs were used for the conciseness model and 100 epochs for clarity. The model is actually not fully converged at these epoch levels, but we chose to limit the training epochs in our submission for the sake of performance, since full convergence can take many hours of training.

A final interesting note is that our conciseness model performed slightly better with un-normalized data, while the clarity model performed significantly better with normalized data.

Model Experiment Design and Evaluation

To evaluate our model, we utilized the StratifiedKFold class from the scikit-learn library to perform a 5-fold cross validation. The data is split into 5 stratified parts each representative

of the original data, then the model is trained on four parts, tested on the remaining, and the resulting RMSE scores are averaged for a final score.

Using our final model parameters and training parameters, we received a cross validation score of **0.352** for conciseness and **0.217** for clarity.

Model Analysis and Challenges

Our deep neural network performed the best out of all the learning models we tried, and contains the most potential for improved performance with more refinement. The biggest challenge and drawback of our DNN model was the extremely long training time required to test and analyze the model as we attempted to tune hyperparameters. After factoring in cross-validation times, the model could easily take several hours to train both the clarity and conciseness models. This made it very difficult to systematically test hyperparameters to determine the optimal settings. With substantially more processing power and time it would have been possible to systematically refine the model with an organized random search for optimal hyperparameters.

For future improvements, the most important factor would be improved feature generation and selection. We saw the largest improvements to our neural networks improvements when word embeddings using the GloVe model were added to the feature set. Further tuning of the model only resulted in relatively small performance increases. We also note that at the level of training selected we had yet to observe any evidence of model overfitting, suggesting that if more computational power and time was available for training, the model could further converge.

Related work

Other participants who participated in this challenge have used a variety of different approaches. One such strategy was to project tree-based models like xgboost, lightgbm, random forest and adaboost, then combining such models to perform ensembling, and then further using outputs of a RNN to account for sentiment meaning to build a complex multi-ensemble model. Another team utilized another ensemble based bagging model combining extreme trees, random forest, stochastic gradient descent, logistic regression, ridge regression, naive bayes, xgboost, light gradient boosting, and word2vec. These models were combined using the stacked ensemble method after selecting the best models.

We definitely took some inspiration while training our models from these papers, including trying out Word2Vec for feature generation.

1. Hong Diao WEN. 2017. CIKM AnalytiCup 2017 – Lazada Product Title Quality Challenge. In Proceedings of ACM CIKM conference, Singapore, October 2017 (CIKM'17), 4 pages. https://doi.org/10.475/123 4

2. Tam T. Nguyen, Hossein Fani, Ebrahim Bagheri, and Gilberto Titericz. 2017. Bagging Model for Product Title Quality with Noise. In Proceedings of International Conference on Information and Knowledge Management, Pan Pacific Singapore, 6-10 Nov. 2017 (CIKM'17), 4 pages. https://doi.org/10.1145/nnnnnnn.nnnnnnn

Conclusion

As an overarching conclusion to the project, it's hard to call any method as superior all around. If anything, it reveals a case of trade-offs between time needed and accuracy desired. Each method had their pros and cons. The three implementations of Naïve Bayes that were tested were relatively easy to set up, very fast to train, and very fast to give results. It wasn't the best when it comes to how accurate it was, returning the highest RMSE scores out of the methods tested, but it was certainly the fastest. Training time using the same amount of training data took less than a minute, and prediction time was equally fast. The best concise score was 0.479, while the best clarity score was 0.237. These where, however, outliers to the various other scores received from different parameters used.

SVM played a down-the-middle line out of the three methods tested. With an RMSE of 0.435 for conciseness, and 0.230 for clarity, its concise score exceeds that of any of the Naïve Bayes scores, while its clarity score outperformed Naïve Bayes' averaged scores from all its tests. The better scores does come with the price of a longer running time, however. Training for the SVM model took hours to complete, a significant increase from the mere minute used by the Naïve Bayes methods. It should be noted that faster methods to train the SVM model exist and could be investigated for use in future works. Traditional SVMs are better suited for smaller datasets due to their longer running time.

Lastly, the Deep Neural Network (DNN) method gave the best scores out of the methods tried, at 0.352 for conciseness and 0.217 for clarity, beating both the scores from running SVM, and the average scores for Naïve Bayes. Training and tuning the model were, like in the SVM case, in the scale of hours. The long training time means tuning was also slow as it required time to complete training before tuning can be done. DNN does have an advantage where it could be parallelized easier, and in theory could allow for better performances and scaling. Overall, while SVM and DNN gave better results, they each have drawbacks. Naïve Bayes may have been the worse in accuracy, but it was also the fastest to work it. It's hard to give a solid winner for each the methods, as under different operational requirements, each one of them has a place. Naïve Bayes can do its calculations much faster as a tradeoff for being less accurate. SVM gives better results on smaller datasets in exchange for longer training times. DNN gives the best results while giving up time needed and resource requirements on the system it's running on.

References

1. Scikit-learn developers. 1.9. Naïve Bayes. Retrieved March 16, 2018 from http://scikit-learn.org/stable/modules/naive_bayes.html

2. Scikit-learn developers. GaussianNB. Retrieved March 16, 2018 from http://scikit-learn.org/stable/modules/generated/sklearn.naive_bayes.GaussianNB.html

3. Scikit-learn developers. MultinomialNB. Retrieved March 16, 2018 from http://scikit-learn.org/stable/modules/generated/sklearn.naive_bayes.MultinomialNB.html

4. Scikit-learn developers. BernoulliNB. Retrieved March 16, 2018 from http://scikit-learn.org/stable/modules/generated/sklearn.naive_bayes.BernoulliNB.html

5. SVM RBF Parameters.
http://scikit-learn.org/stable/auto_examples/svm/plot_rbf_parameters.html

6. How to implement SMOTE in GridSearchCV.
https://stackoverflow.com/questions/48370150

7. Which search range for determining optimal C and gamma parameters.
https://stats.stackexchange.com/questions/43943

8. SciKit Learn SVM.
http://scikit-learn.org/stable/modules/svm.html

9. Gensim Doc2Vec.
https://radimrehurek.com/gensim/models/doc2vec.html

10. SciKit Learn CountVectorizer.
http://scikit-learn.org/stable/modules/generated/sklearn.feature_extraction.text.CountVectorizer.html

11. BeautifulSoup 4.6.0,
https://www.crummy.com/software/BeautifulSoup/

12. Duchi, John, Elad Hazan, and Yoram Singer. "Adaptive subgradient methods for online learning and stochastic optimization." Journal of Machine Learning Research 12.Jul (2011): 2121-2159.

13. Krizhevsky, Alex, et al. "ImageNet Classification with Deep Convolutional Neural Networks."Communications of the ACM, vol. 60, no. 6, 2017, pp. 84–90., doi:10.1145/3065386.

14. Pennington, Jeffrey, Richard Socher, and Christopher Manning. "Glove: Global vectors for word representation." Proceedings of the 2014 conference on empirical methods in natural language processing (EMNLP). 2014.

15. Spacy 2.0.9, "Industrial-Strength Natural Language Processing", https://spacy.io/

16. Srivastava, Nitish, et al. "Dropout: A simple way to prevent neural networks from overfitting." The Journal of Machine Learning Research 15.1 (2014): 1929-1958.

Task and work load distribution

- Rodrigo Valle - SVMs, writer, editor, typesetting
- Krit Saefang - Naive Bayes, writer, editor
- Bill Tang - Neural networks, packager

Third Party Packages Used

- scikit-learn
- imblearn
- numpy
- pandas

- gensim
- tensorflow
- keras
- spacy
- beautifulsoup4